

CPME 452 Assignment 2 – Backpropagation

Adam Grace

20055169

October 25th, 2020

Model 1 – From Scratch

I create a 784 input, 10 output backpropagation neural network using momentum to classify the handwritten digits from the MNIST dataset. In training and in testing I achieved an overall accuracy of 90.95% and 91.13% respectively. The parameters that I varied while investigating this model include the number nodes in the hidden layer, the number of epochs, the learning rate, the momentum rate, and the initial weight range. After trying a limited number of variable combinations, I settled on 30 hidden nodes, 5 epochs, a learning rate of 0.01, a momentum rate of 0.3 and an initial weight range of negative 1 to 1. Initial weights were chosen randomly from this range. Unfortunately, I was not able to change around the variables more than about 10 times due to the long computation time. For data preprocessing, I flattened the 28x28 2D input array into a 784 element 1D array. I also found that a higher overall accuracy was achieved when I introduced input normalization. The confusion matrices and the precision, recall and overall accuracy statistics are shown below. The confusion matrices goes from 0 to 9, starting from the top left.

TRAINING PHASE

Overall accuracy:

90.9483333333334 %

	precision	recall	f1-score	support
0	0.93	0.97	0.95	5923
1	0.94	0.97	0.96	6742
2	0.92	0.87	0.89	5958
3	0.91	0.88	0.89	6131
4	0.90	0.91	0.91	5842
5	0.88	0.86	0.87	5421
6	0.92	0.95	0.94	5918
7	0.93	0.92	0.93	6265
8	0.87	0.87	0.87	5851
9	0.88	0.89	0.88	5949
accuracy			0.91	60000
macro avg	0.91	0.91	0.91	60000
weighted avg	0.91	0.91	0.91	60000

```
[[5729 1 17 14 7 43 44 4 57 7]
 [ 2 6545 30 21 11 33 14 18 47 21]
 [ 72 70 5194 108 96 21 125 105 130 37]
 [ 45 27 135 5385 14 209 34 64 152 66]
 [ 17 37 23 1 5326 6 83 15 60 274]
 [ 99 25 36 172 93 4682 113 29 116 56]
 [ 57 23 32 1 32 85 5628 6 50 4]
 [ 33 66 86 26 92 15 2 5749 23 173]
 [ 46 119 75 129 46 177 52 32 5062 113]
 [ 51 28 26 85 199 43 11 143 94 5269]]
```

TESTING PHASE

Overall accuracy:

91.13 %

	precision	recall	f1-score	support
0	0.93	0.98	0.95	980
1	0.96	0.98	0.97	1135
2	0.94	0.88	0.91	1032
3	0.90	0.89	0.90	1010
4	0.90	0.91	0.91	982
5	0.87	0.87	0.87	892
6	0.91	0.94	0.92	958
7	0.93	0.90	0.92	1028
8	0.87	0.86	0.87	974
9	0.89	0.89	0.89	1009
accuracy			0.91	10000
macro avg	0.91	0.91	0.91	10000
weighted avg	0.91	0.91	0.91	10000

```
[[ 958 0 1 2 0 4 9 1 5 0]
 [ 0 1111 3 2 1 1 3 2 12 0]
 [ 14 10 913 18 17 5 12 17 21 5]
 [ 5 1 15 901 2 35 3 12 25 11]
 [ 3 3 1 1 894 2 21 2 11 44]
 [ 17 2 5 26 13 778 20 6 17 8]
 [ 15 3 3 2 14 17 898 1 5 0]
 [ 3 19 23 8 9 2 1 928 6 29]
 [ 10 7 6 26 7 34 17 11 837 19]
 [ 7 7 1 15 32 13 5 15 19 895]]
```

Model 2A – Using Predefined Libraries and Parameters from Model 1

Using the same parameters as in Model 1, I created the same neural network using predefined libraries from Keras. In training and in testing I achieved an overall accuracy of 89.98% and 90.60% respectively. The confusion matrices and the precision, recall and overall accuracy statistics are shown below. The confusion matrices goes from 0 to 9, starting from the top left.

TRAINING PHASE

Overall accuracy:

89.975 %

	precision	recall	f1-score	support
0	0.93	0.97	0.95	5923
1	0.92	0.97	0.94	6742
2	0.90	0.87	0.88	5958
3	0.90	0.86	0.88	6131
4	0.89	0.91	0.90	5842
5	0.89	0.80	0.84	5421
6	0.91	0.95	0.93	5918
7	0.92	0.91	0.92	6265
8	0.88	0.86	0.87	5851
9	0.86	0.88	0.87	5949
accuracy			0.90	60000
macro avg	0.90	0.90	0.90	60000
weighted avg	0.90	0.90	0.90	60000

```
[[5728  0  29  13  19  20  46  5  58  5]
 [  0 6536 39  22  5  37  13 19  58 13]
 [ 69  84 5183 78 119  13 125 124 134 29]
 [ 36  43 180 5283 8 229  53  87 118 94]
 [  8  48  35  2 5310  4  83  14  39 299]
 [137  78  58 231 105 4346 144 35 203 84]
 [ 58  24  49  4  40  69 5642  0  31  1]
 [ 39 109  81  15  82  9  3 5730 13 184]
 [ 49 177  85 150  41 137  67 29 5006 110]
 [ 62  38  33 104 214  38  6 187  46 5221]]
```

TESTING PHASE

Overall accuracy:

90.6 %

	precision	recall	f1-score	support
0	0.92	0.98	0.95	980
1	0.95	0.98	0.96	1135
2	0.92	0.87	0.89	1032
3	0.90	0.89	0.90	1010
4	0.89	0.91	0.90	982
5	0.91	0.81	0.86	892
6	0.91	0.95	0.93	958
7	0.91	0.90	0.91	1028
8	0.88	0.86	0.87	974
9	0.87	0.88	0.88	1009
accuracy			0.91	10000
macro avg	0.91	0.90	0.90	10000
weighted avg	0.91	0.91	0.91	10000

```
[[ 963  0  2  1  0  0  9  1  4  0]
 [  0 1108  1  5  1  1  5  2 12  0]
 [ 14  7 901 16 17  0 17 21 33  6]
 [  4  2 21 902  1 27  5 17 18 13]
 [  1  5  5  0 898  1 14  2  5 51]
 [ 20  4  7 45 16 723 19 13 37  8]
 [ 16  3  4  1  8 12 910  1  3  0]
 [  6 22 27  2 11  0  0 926  1 33]
 [ 10 12 11 22  9 21 18 14 841 16]
 [ 18  8  3 11 45  9  2 18  7 888]]
```

Model 2B – Using Predefined Libraries and New Parameters

Model 2A was used with new parameters in an attempt to achieve better results. 100 hidden nodes, 5 epochs, a learning rate of 0.1, a momentum rate of 0.5 and an initial weight range of negative 1 to 1 was used. In training and in testing I improved upon previous models, achieving an overall accuracy of 89.98% and 90.60% respectively. The confusion matrices and the precision, recall and overall accuracy statistics are shown below. The confusion matrices goes from 0 to 9, starting from the top left.

TRAINING PHASE

Overall accuracy:

96.29 %

	precision	recall	f1-score	support
0	0.97	0.98	0.98	5923
1	0.98	0.98	0.98	6742
2	0.96	0.96	0.96	5958
3	0.95	0.95	0.95	6131
4	0.94	0.97	0.95	5842
5	0.97	0.95	0.96	5421
6	0.98	0.97	0.98	5918
7	0.97	0.97	0.97	6265
8	0.94	0.97	0.95	5851
9	0.97	0.92	0.95	5949
accuracy			0.96	60000
macro avg	0.96	0.96	0.96	60000
weighted avg	0.96	0.96	0.96	60000

```
[[5818 1 7 5 9 9 24 5 41 4]
 [ 1 6581 37 36 15 11 2 13 40 6]
 [ 25 15 5701 42 53 3 11 50 52 6]
 [ 3 10 65 5850 5 53 14 46 66 19]
 [ 13 10 23 3 5677 2 21 9 19 65]
 [ 25 6 17 82 25 5171 32 6 36 21]
 [ 37 10 8 3 27 56 5755 0 22 0]
 [ 12 20 40 20 44 5 2 6078 14 30]
 [ 19 35 21 47 17 26 18 7 5651 10]
 [ 28 10 12 82 181 17 1 68 58 5492]]
```

TESTING PHASE

Overall accuracy:

95.86 %

	precision	recall	f1-score	support
0	0.95	0.99	0.97	980
1	0.98	0.98	0.98	1135
2	0.96	0.96	0.96	1032
3	0.94	0.97	0.95	1010
4	0.94	0.97	0.95	982
5	0.96	0.93	0.95	892
6	0.96	0.96	0.96	958
7	0.96	0.96	0.96	1028
8	0.95	0.97	0.96	974
9	0.97	0.91	0.94	1009
accuracy			0.96	10000
macro avg	0.96	0.96	0.96	10000
weighted avg	0.96	0.96	0.96	10000

```
[[ 968 0 0 4 0 3 2 1 1 1]
 [ 0 1115 2 3 0 1 5 2 7 0]
 [ 8 0 988 7 8 0 4 7 10 0]
 [ 0 0 10 978 0 6 1 9 6 0]
 [ 2 1 4 1 949 0 5 3 5 12]
 [ 8 1 1 25 3 827 14 1 8 4]
 [ 10 3 3 2 6 10 917 0 7 0]
 [ 3 10 16 3 4 0 0 983 2 7]
 [ 5 1 2 7 5 4 5 3 942 0]
 [ 10 6 1 13 39 6 0 10 5 919]]
```

Discussion

Interestingly, on all models the testing overall accuracy was greater than the training overall accuracy. This was unexpected since neural networks tend to fit to the data it is trained on. This indicates that overfitting was not an issue. In Model 2B, I achieved much better overall accuracy, precision and recall. I believe that the greater part of this improvement is due to the increase in hidden nodes from 30 to 100. It was observed that models 1 and 2A had some difficulty classifying 5 and 9, possibly due to not enough hidden nodes. Unrelated to the results, I also noticed that the predefined libraries were able to train the network in less time than my program from scratch, likely due to operation optimization.